


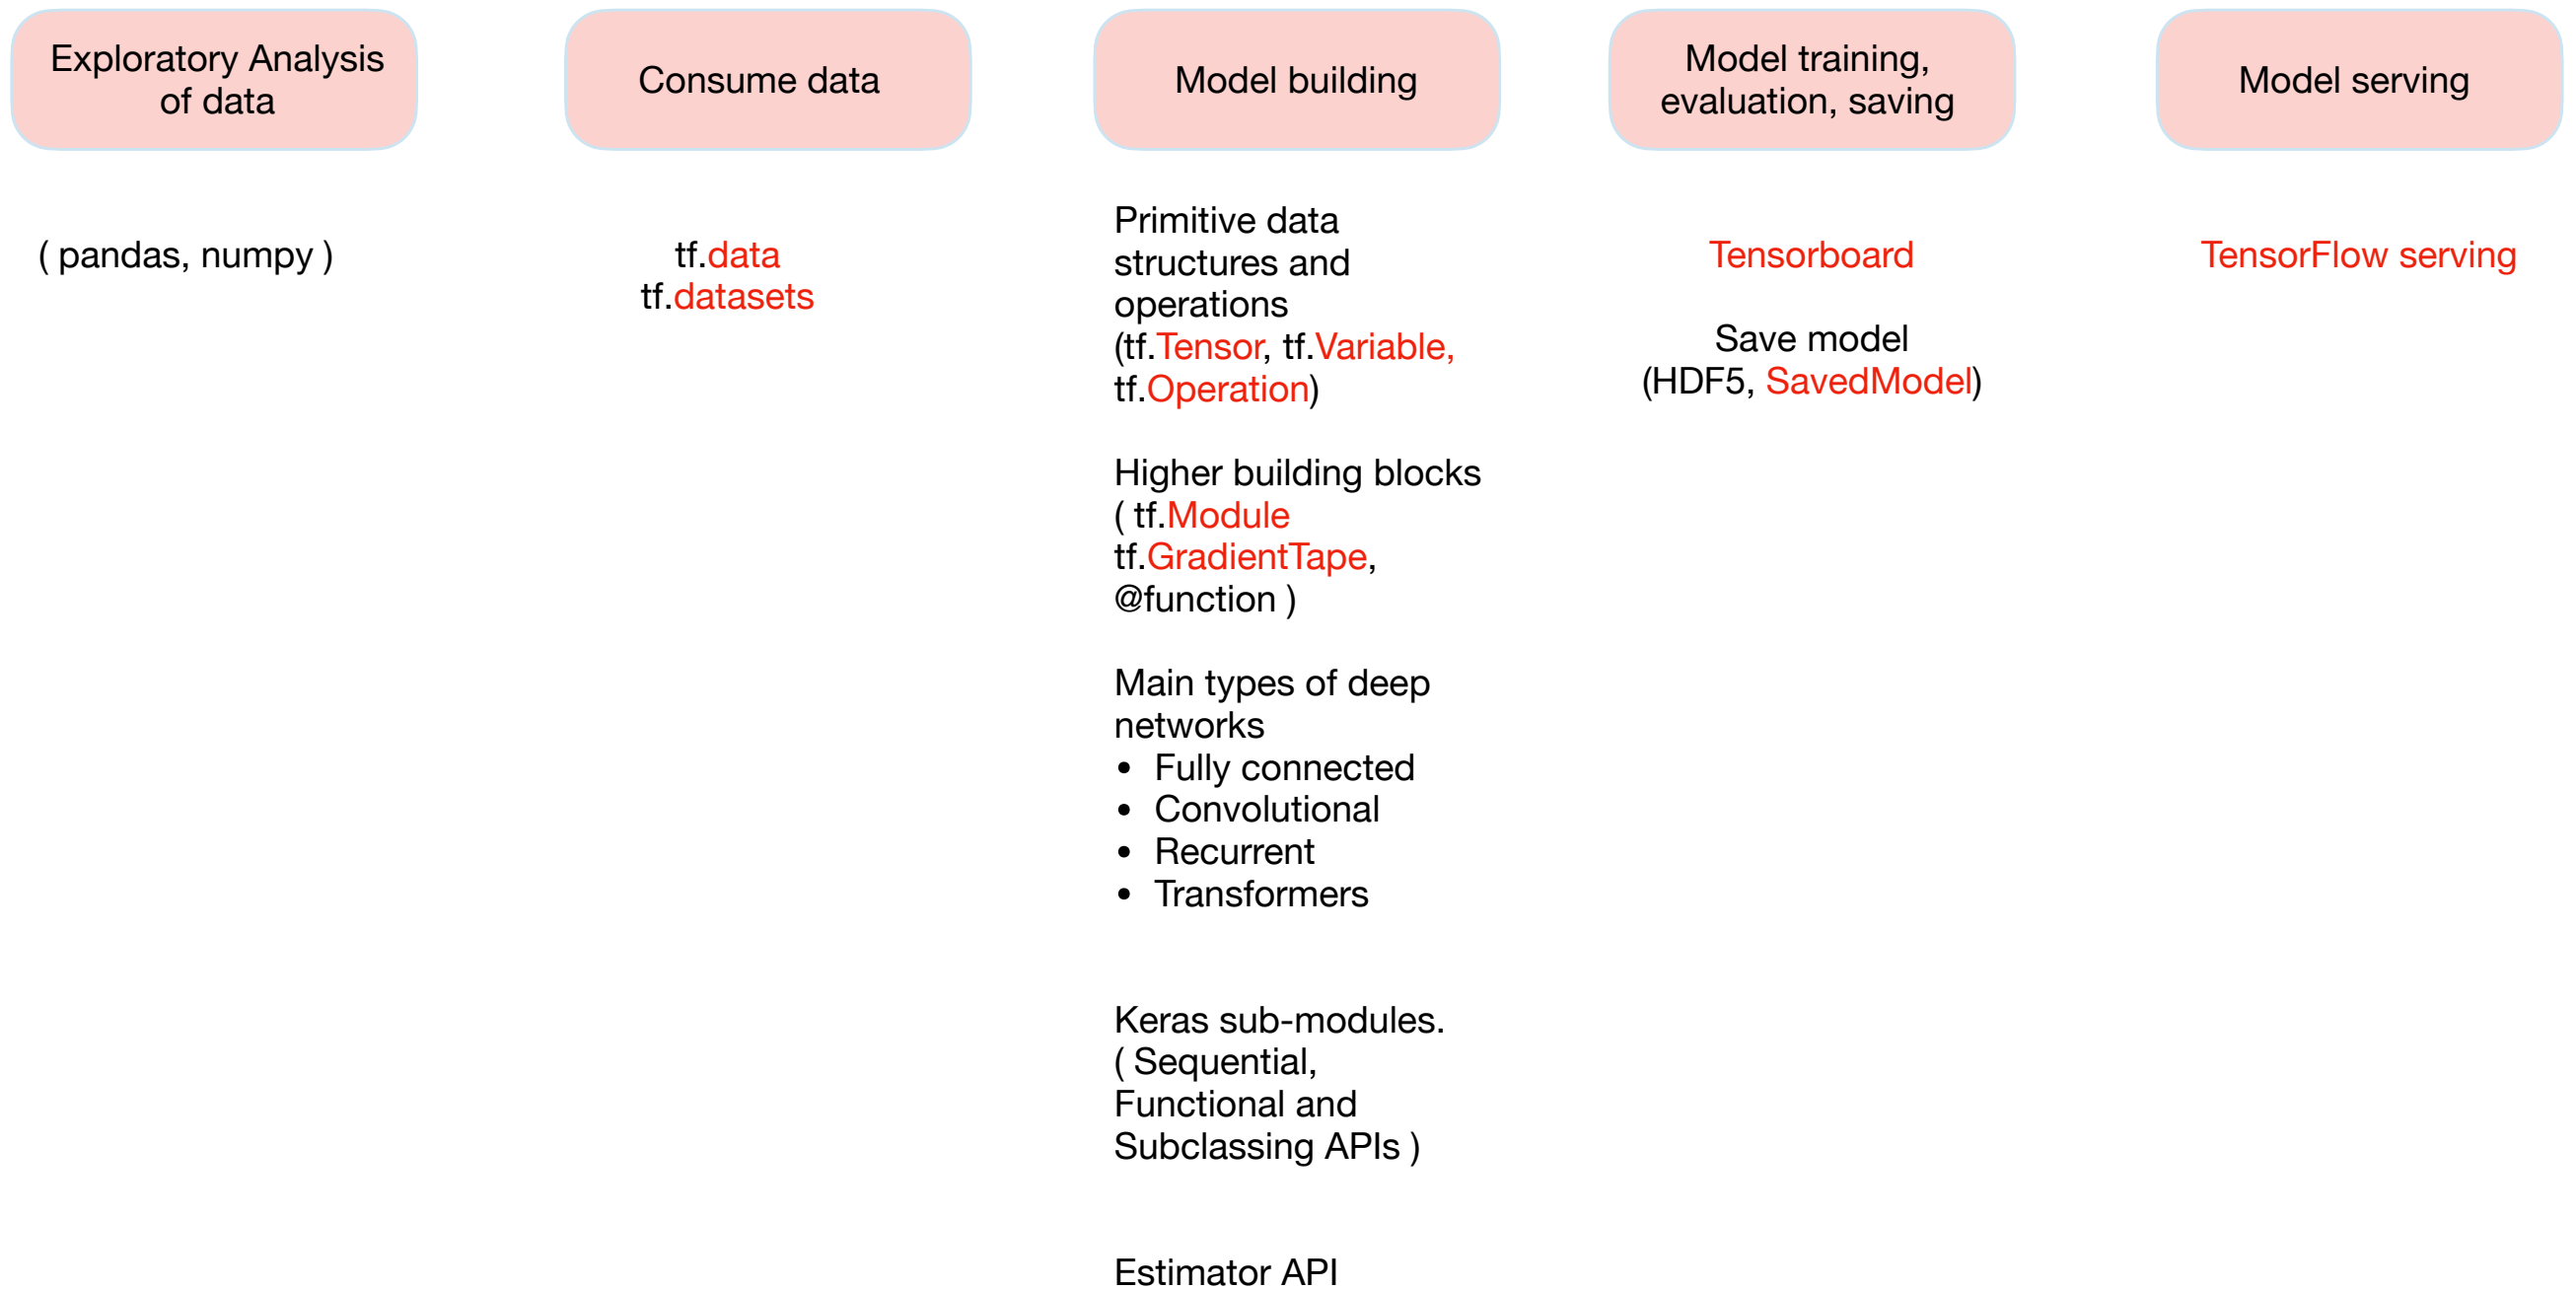
TensorFlow 2.0



Machine Learning



Overview of TensorFlow components



TensorFlow

building-blocks



1- tf.Tensor

```
>>> import tensorflow as tf
>>> tf.__version__
```

tf.Tensor is an n-dimensional data structure heavily used when defining machine learning models, as they are used to store inputs, intermediate outputs of layers and final outputs of the model. **tf.Tensor** is an immutable object.

```
>>> tf.constant(
    value,
    dtype= None,
    shape=None,
    name= 'Const'
)
>>> a = tf.constant( 2, shape=[4], dtype=tf.float32 )
>>> b = tf.constant( 3, shape=[4], dtype=tf.float32 )
>>> c = tf.add(a, b)
>>> type( c ).__name__
>>> c[0].assign( 1 )
```

```
>>> tf.constant( 0, shape=(2,2) )
>>> tf.constant( [ 1, 2, 3, 4, 5, 6 ] )
>>> tf.constant( np.zeros( (3,3) ) )
```

Create tensor

```
>>> tf.zeros( shape )      tf.zeros_like( input )
>>> tf.ones( shape )      tf.ones_like( input )
>>> tf.eye( num_rows, num_columns=None )
>>> tf.fill( dims, value )
>>> tf.linspace( start, stop, num )
>>> tf.range( start, limit, delta=1 )
```

```
>>> tf.random.normal( shape, mean=0.0, stddev=1.0 )
>>> tf.random.uniform( shape, minval=0, maxval=None )
```

Like NumPy arrays [start : end : step]

Indexing

```
>>> t[ :, 0 ], t[ t < 2 ], t[ 2 : 7 ].numpy()      tf.gather( t, indices=tf.constant([0,2]) )
```

```
>>> t.device
>>> t.dtype
>>> t.name
>>> t.shape
>>> t.ndim
```

Attributes and Methods

```
>>> tf.cast( x, dtype )
>>> tf.reshape( tensor, shape )
>>> tf.transpose( a, perm=None )
>>> tf.where( condition, x=None, y=None )
>>> tf.squeeze( input, axis=None ) tf.expand_dims( input, axis=None )
>>> tf.sort( values, axis=-1, direction='ASCENDING' ) tf.argsort( ... )
>>> tf.stack( values, axis=0 ) tf.unstack( values, axis=0 )
>>> tf.one_hot( indices, depth )
```

Common operations

```
>>> tf.add( x, y )      x + y
>>> tf.subtract( x, y )  x - y
>>> tf.multiply( x, y )  x * y
>>> tf.divide( x, y )    x / y
```

```
>>> tf.matmul( A, B )    A @ B
>>> tf.linalg.matvec( A, x )
>>> tf.norm( t , ord='euclidean')
```

Math

```
>>> tf.abs( x )
>>> tf.math.sigmoid( x )
>>> tf.math.tanh( x )
>>> tf.math.cos( x ) tf.math.acos( x )
>>> tf.math.exp( x ) tf.math.log( x )
>>> tf.math.pow( x, y ) x**2
>>> tf.math.floor( x ) tf.math.ceil( x )
>>> tf.math.argmax( input, axis=None )
>>> tf.math.argmin( input, axis=None )
>>> tf.math.top_k( input, k=1, sorted=True )
```

```
>>> tf.reduce_sum( input_tensor, axis=None )
>>> tf.reduce_prod( input_tensor, axis=None )
>>> tf.reduce_min( input_tensor, axis=None )
>>> tf.reduce_max( input_tensor, axis=None )
>>> tf.math.cumsum( x, axis=0 )
>>> tf.math.cumprod( x, axis=0 )
```

2- tf.Variable

tf.Variable is ideal for defining model parameters and it can change the value of elements as required after it is initialized. **tf.Variable** is a mutable object.

```
>>> tf.Variable(
    initial_value= [ tensor, numpy, initializer ],
    trainable= True,
    dtype= None,
    name= None
)
```

Create a variable

tf.keras.initializers

```
Constant( value )
Ones(), Zeros(), Identity()
GlorotNormal(), GlorotUniform()
HeNormal(), HeUniform()
RandomNormal(), RandomUniform()
```

dtype

```
tf.float16, 32, 64
tf.uint8, 16, 32, 64
tf.int8, 16, 32, 64
tf.bool
```

```
>>> v1 = tf.Variable( tf.constant(2.0, shape=[4]), dtype=tf.float32 )
>>> v2 = tf.Variable( np.ones(shape=[4,3]), dtype='float32' )
>>> v3 = tf.Variable( tf.keras.initializers.RandomNormal() ( shape=[3,4,5] ) )
```

```
>>> v.device
>>> v.dtype
>>> v.name
>>> v.shape
>>> v.trainable
>>> v.numpy()
>>> v.assign()
>>> v = tf.Variable( np.zeros(shape=[4,3] ) )
>>> v = v[0,2].assign(1)
>>> v[2:, :2].assign( [[3,3],[3,3]] )
```

Attributes and Methods



TensorFlow Concepts

```
import tensorflow as tf
import matplotlib.pyplot as plt

# -----Generate data ----- #
TRUE_W = 3.0
TRUE_B = 2.0
NUM_EXAMPLES = 1000

x = tf.random.normal( shape=[NUM_EXAMPLES] )
noise = tf.random.normal( shape=[NUM_EXAMPLES] )
y = x * TRUE_W + TRUE_B + noise

# -----Define model ----- #
class MyModel( tf.Module ):
    def __init__(self, **kwargs):
        super().__init__( **kwargs )
        self.w = tf.Variable(5.0)
        self.b = tf.Variable(0.0)

    def __call__( self, x ):
        return self.w * x + self.b

# -----Define loss ----- #
@tf.function
def loss(target_y, predicted_y):
    return tf.reduce_mean( t f.square(target_y - predicted_y) )

# -----Train a model ----- #
model = MyModel()
learning_rate= 0.1
epochs = 25

for epoch in range(epochs):
    with tf.GradientTape( ) as tape:
        yhat = model( x )
        current_loss = loss(y, yhat)

    dw, db = tape.gradient( current_loss, [model.w, model.b] )
    model.w.assign_sub( learning_rate * dw )
    model.b.assign_sub( learning_rate * db )

    print( "Epoch %2d: loss=%2.5f" % (epoch, current_loss) )

plt.scatter( x, y, c="b" )
plt.plot(x, model(x), c="r")
plt.show( )
```

Tensors, `tf.Tensor`, are multi-dimensional arrays with a uniform type. A `tf.Variable` represents a tensor whose value can be changed by running ops on it.

Automatic differentiation is useful for implementing machine learning algorithms such as backpropagation for training neural networks.

TensorFlow provides the `tf.GradientTape` API for automatic differentiation. TensorFlow "records" relevant operations executed inside the context of a `tf.GradientTape` onto a "tape". TensorFlow then uses that tape to compute the gradients of a "recorded" computation using reverse mode differentiation.

Graphs are data structures that contain a set of `tf.Operation` objects, which represent units of computation; and `tf.Tensor` objects, which represent the units of data that flow between operations. They are defined in a `tf.Graph` context; graphs are extremely useful and let your TensorFlow run fast, run in parallel, and run efficiently on multiple devices.

You create and run a graph in TensorFlow by using `tf.function`, either as a direct call or as a decorator. `tf.function` takes a regular function as input and returns a **Function**. A Function is a Python callable that builds TensorFlow graphs from the Python function. You use a Function in the same way as its Python equivalent.

While TensorFlow operations are easily captured by a `tf.Graph`, Python-specific logic needs to undergo an extra step in order to become part of the graph. `tf.function` uses a library called **AutoGraph** (`tf.autograph`) to convert Python code into graph-generating code.


Most **models** are made of layers. Layers are functions with a known mathematical structure that can be reused and have trainable variables. In TensorFlow, most high-level implementations of **layers** and **models**, such as Keras, are built on the same foundational class: `tf.Module`. By **subclassing** `tf.Module`, any `tf.Variable` or `tf.Module` instances assigned to this object's properties are **automatically** and **recursively** collected.

tf.keras




1- tf.keras

Used to implement networks that are linear and has one input and one output



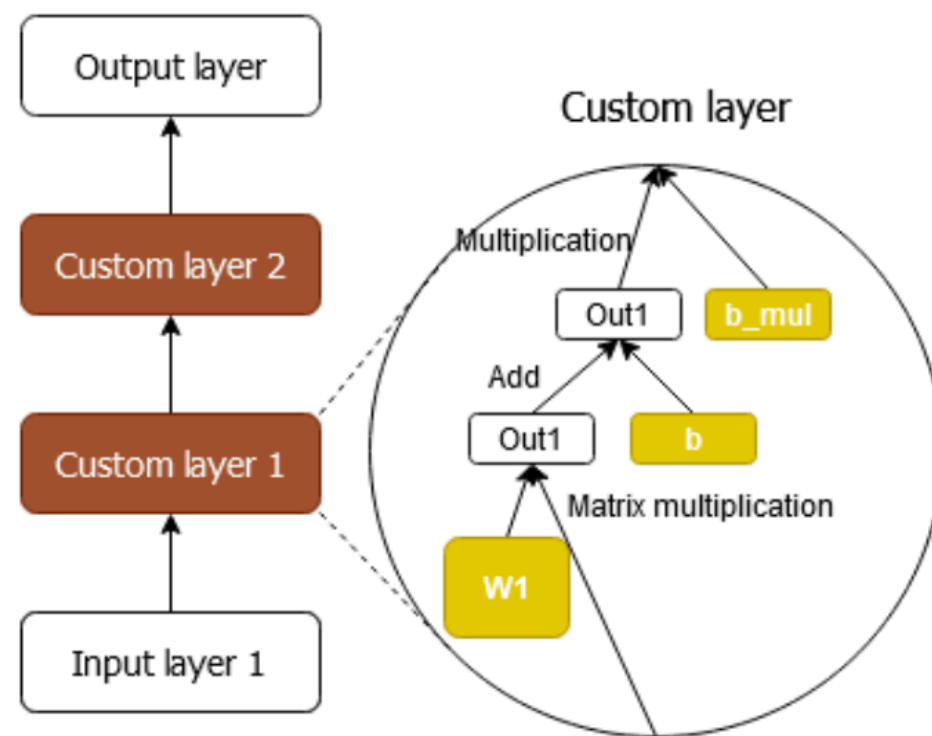
Sequential API

Allows more flexibility such as having multiple input layers, parallel hidden layers and multiple outputs



Functional API

Enables to implement highly custom layer computations and models by implementing a layer/model as a Python object



Sub-classing API

Flexibility

Ease of using



Keras: standard training

```
import tensorflow as tf

# -----Define model ----- #
model = tf.keras.models.Sequential( )

model.add( tf.keras.layers.Flatten( input_shape=x_train[0].shape) )

model.add( tf.keras.layers.Dense( 128, activation='tanh') )
model.add( tf.keras.layers.Dense( 128, activation='tanh') )
model.add( tf.keras.layers.Dropout( 0.1 ) )
model.add( tf.keras.layers.Dense(10, activation='softmax') )

# -----Compile model ----- #
model.compile(
    optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

# -----Fit model ----- #
result = model.fit(
    x_train,
    y_train,
    validation_data=(x_test, y_test),
    epochs=10)
```

Model: Attributes and Methods

```
>>> model.inputs
>>> model.outputs
>>> model.trainable_variables
>>> model.layers

>>> model.get_layer(name)
```

Layer: Attributes and Methods

```
>>> layer.weights
>>> layer.input
>>> layer.input_shape
>>> layer.output
>>> layer.output_shape
>>> layer.trainable
```

```
feature_extractor = tf.keras.Model(
    inputs = initial_model.inputs,
    outputs = [ layer.output for layer in initial_model.layers ]
)

features = feature_extractor( x )
```

Linear Layers

Linear Layers

```
# Dense implements the operation: output = activation( dot( input, kernel ) + bias )
>>> tf.keras.layers.Dense( units, activation=None, [ input_shape, name ] )
```

Dropout Layers

```
# The Dropout layer randomly sets input units to 0 with a frequency of rate at each step
during training time
>>> tf.keras.layers.Dropout( rate, [ input_shape, name ] )
```

Utility Functions

```
# Input() is used to instantiate a Keras tensor.
>>> tf.keras.Input( shape )

# Flattens the input.
>>> tf.keras.layers.Flatten( [ input_shape, name ] )

# Layer that reshapes inputs into the given shape
>>> tf.keras.layers.Reshape( target_shape, [ input_shape, name ] )


# Layer that concatenates a list of inputs.
>>> tf.keras.layers.Concatenate( axis=-1 [input_shape, name ] )
```




Activations


```
# [0,1]
>>> tf.keras.activations.sigmoid()
```

$$\text{Sigmoid}(x) = \sigma(x) = \frac{1}{1 + \exp(-x)}$$




```
# [-1, 1]
>>> tf.keras.activations.tanh()
```

$$\text{Tanh}(x) = \tanh(x) = \frac{\exp(x) - \exp(-x)}{\exp(x) + \exp(-x)}$$



```
# [0, +inf]
>>> tf.keras.activations.relu()
```

$$\text{ReLU}(x) = (x)^+ = \max(0, x)$$



```
# [0, 1]
>>> tf.keras.activations.softmax()
```

$$\text{Softmax}(x_i) = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$

Losses

Regression

```
# mean of squares of errors between labels and predictions.
# model.compile( optimizer='sgd', loss=tf.keras.losses.MeanSquaredError() )
>>> tf.keras.losses.MeanSquaredError()
```

```
# mean of absolute difference between labels and predictions
# model.compile( optimizer='sgd', loss=tf.keras.losses.MeanAbsoluteError() )
>>> tf.keras.losses.MeanAbsoluteError()
```

Classification

```
# Use this cross-entropy loss for binary (0 or 1) classification applications
# model.compile( loss=tf.keras.losses.BinaryCrossentropy(from_logits=True) )
>>> tf.keras.losses.BinaryCrossentropy( from_logits=False )
```

```
# Use this crossentropy loss function when there are two or more label classes.
# We expect labels to be provided in a one_hot representation.
# model.compile( loss=tf.keras.losses.CategoricalCrossentropy() )
>>> tf.keras.losses.CategoricalCrossentropy( from_logits=False )
```

```
# Use this crossentropy loss function when there are two or more label classes.
# We expect labels to be provided as integers.
# model.compile( loss=tf.keras.losses.SparseCategoricalCrossentropy() )
>>> tf.keras.losses.SparseCategoricalCrossentropy( from_logits=False )
```

Optimization Algorithms

```
# Gradient descent (with momentum) optimizer
>>> tf.keras.optimizers.SGD( learning_rate=0.01,
                             momentum=0.0,
                             nesterov=False )
```

```
# Optimizer that implements the Adam algorithm
>>> tf.keras.optimizers.Adam( learning_rate=0.001,
                              beta_1=0.9,
                              beta_2=0.999,
                              eps=1e-08,
                              amsgrad=False )
```